

## Keras -- MLPs on MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 55s 5us/step

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point
print(X_train[0])
```

[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	82	
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154		
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0															

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normalizing
print(X_train[0])
```

[0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0	0	0	0	0	0

[illegible]



```
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal
```

In [12]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## 1. 2-hidden layer architecture (784-352-124-10)

### 1.1 MLP + ReLU activation + ADAM Optimizer

In [15]:

```
model_relu = Sequential()
model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(
seed=None)))
```

```

model_relu.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:4185: The name tf.truncated\_normal is deprecated. Please use tf.random.truncated\_normal instead.

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:4138: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 352)	276320
dense_2 (Dense)	(None, 124)	43772
dense_3 (Dense)	(None, 10)	1250

=====  
 Total params: 321,342  
 Trainable params: 321,342  
 Non-trainable params: 0

None

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\tensorflow\python\ops\math\_grad.py:1250: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 87us/step - loss: 0.2540 - acc: 0.9269 - val\_loss: 0.1213 - val\_acc: 0.9621

Epoch 2/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0975 - acc: 0.9714 - val\_loss: 0.0903 - val\_acc: 0.9711

Epoch 3/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0621 - acc: 0.9803 - val\_loss: 0.0731 - val\_acc: 0.9783

Epoch 4/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0429 - acc: 0.9869 - val\_loss: 0.0655 - val\_acc: 0.9788

Epoch 5/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0318 - acc: 0.9898 - val\_loss: 0.0664 - val\_acc: 0.9782

Epoch 6/20

60000/60000 [=====] - 5s 76us/step - loss: 0.0231 - acc: 0.9928 - val\_loss: 0.0749 - val\_acc: 0.9789

Epoch 7/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0191 - acc: 0.9943 - val\_loss: 0.0820 - val\_acc: 0.9769

Epoch 8/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0187 - acc: 0.9941 - val\_loss: 0.0753 - val\_acc: 0.9783

Epoch 9/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0139 - acc: 0.9953 - val\_loss: 0.0653 - val\_acc: 0.9806

Epoch 10/20

```

Epoch 10/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0099 - acc: 0.9970 -
val_loss: 0.0779 - val_acc: 0.9804
Epoch 11/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0121 - acc: 0.9960 -
val_loss: 0.0783 - val_acc: 0.9800
Epoch 12/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0111 - acc: 0.9968 -
val_loss: 0.0734 - val_acc: 0.9806
Epoch 13/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0078 - acc: 0.9973 -
val_loss: 0.0951 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0075 - acc: 0.9975 -
val_loss: 0.0831 - val_acc: 0.9817
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0069 - acc: 0.9975 -
val_loss: 0.0943 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0127 - acc: 0.9957 -
val_loss: 0.0873 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0053 - acc: 0.9980 -
val_loss: 0.1011 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0083 - acc: 0.9974 -
val_loss: 0.0797 - val_acc: 0.9837
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0057 - acc: 0.9980 -
val_loss: 0.0863 - val_acc: 0.9801
Epoch 20/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0081 - acc: 0.9974 -
val_loss: 0.0876 - val_acc: 0.9816

```

In [16]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

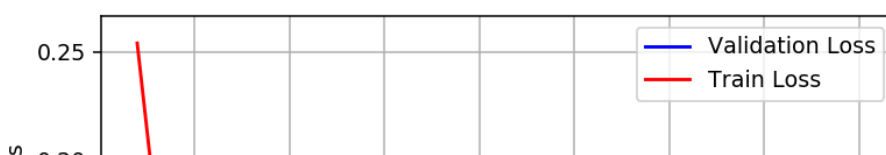
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

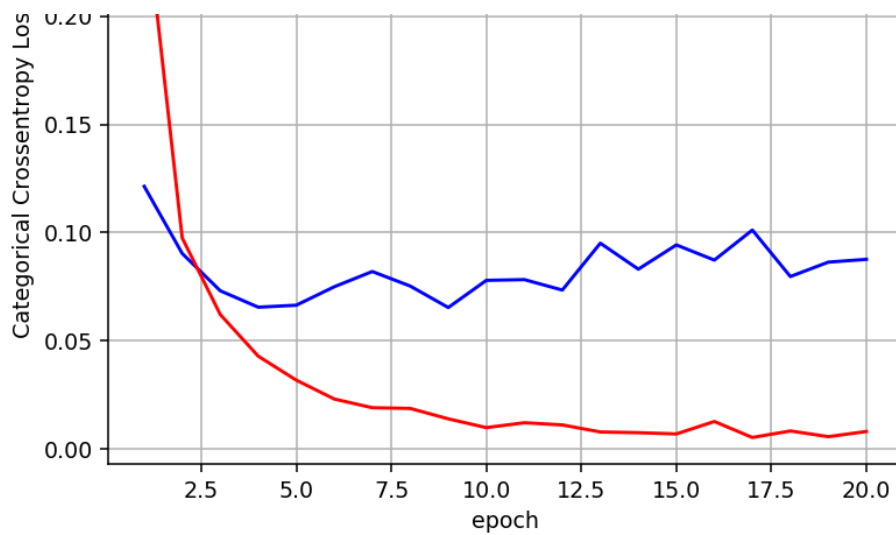
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08760498947395635

Test accuracy: 0.9816





## 1.2 MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [17]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 352)	276320
batch_normalization_1 (Batch Normalization)	(None, 352)	1408
dense_5 (Dense)	(None, 124)	43772
batch_normalization_2 (Batch Normalization)	(None, 124)	496
dense_6 (Dense)	(None, 10)	1250
Total params: 323,246		
Trainable params: 322,294		
Non-trainable params: 952		

In [18]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```



```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 120us/step - loss: 0.2029 - acc: 0.9397 -
val_loss: 0.1044 - val_acc: 0.9672
Epoch 2/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0770 - acc: 0.9766 -
val_loss: 0.0813 - val_acc: 0.9744
Epoch 3/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0479 - acc: 0.9853 -
val_loss: 0.0742 - val_acc: 0.9777
Epoch 4/20
60000/60000 [=====] - 5s 75us/step - loss: 0.0352 - acc: 0.9887 -
val_loss: 0.0810 - val_acc: 0.9763
Epoch 5/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0276 - acc: 0.9907 -
val_loss: 0.0798 - val_acc: 0.9760
Epoch 6/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0219 - acc: 0.9934 -
val_loss: 0.0760 - val_acc: 0.9770
Epoch 7/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0198 - acc: 0.9937 -
val_loss: 0.0728 - val_acc: 0.9815
Epoch 8/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0175 - acc: 0.9944 -
val_loss: 0.0780 - val_acc: 0.9784
Epoch 9/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0132 - acc: 0.9958 -
val_loss: 0.0839 - val_acc: 0.9767
Epoch 10/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0150 - acc: 0.9953 -
val_loss: 0.0710 - val_acc: 0.9810
Epoch 11/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0126 - acc: 0.9954 -
val_loss: 0.0776 - val_acc: 0.9789
Epoch 12/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0124 - acc: 0.9962 -
val_loss: 0.0853 - val_acc: 0.9776
Epoch 13/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0103 - acc: 0.9969 -
val_loss: 0.0754 - val_acc: 0.9807
Epoch 14/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0096 - acc: 0.9968 -
val_loss: 0.0691 - val_acc: 0.9820
Epoch 15/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0103 - acc: 0.9966 -
val_loss: 0.0840 - val_acc: 0.9782
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0100 - acc: 0.9966 -
val_loss: 0.0779 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0095 - acc: 0.9970 -
val_loss: 0.0855 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0079 - acc: 0.9972 -
val_loss: 0.0874 - val_acc: 0.9794
Epoch 19/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0072 - acc: 0.9976 -
val_loss: 0.0891 - val_acc: 0.9800
Epoch 20/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0070 - acc: 0.9978 -
val_loss: 0.0844 - val_acc: 0.9788
```

In [19]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

```

# plt.figure(figsize=(10, 10))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

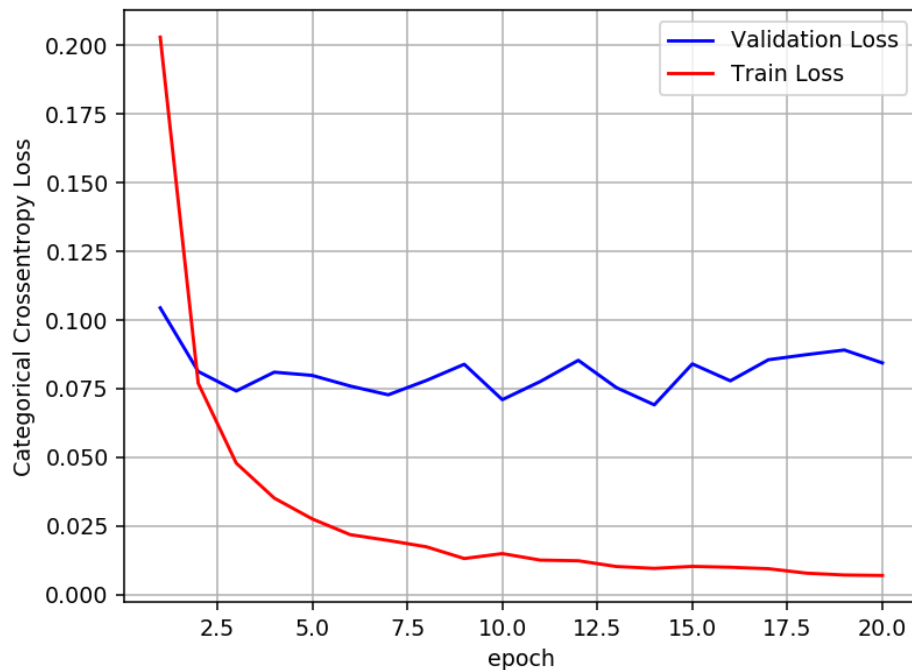
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08438539689127975

Test accuracy: 0.9788



## 1.3 MLP + Dropout + Adam Optimizer

In [20]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

```

```
model_drop.summary()
```

WARNING:tensorflow:From C:\Users\NIKHITHA\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [21]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 9s 145us/step - loss: 0.4204 - acc: 0.8730 - val_loss: 0.1374 - val_acc: 0.9567
Epoch 2/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2043 - acc: 0.9384 - val_loss: 0.1108 - val_acc: 0.9642
Epoch 3/20
60000/60000 [=====] - 8s 132us/step - loss: 0.1579 - acc: 0.9527 - val_loss: 0.0897 - val_acc: 0.9726
Epoch 4/20
60000/60000 [=====] - 8s 132us/step - loss: 0.1358 - acc: 0.9590 - val_loss: 0.0782 - val_acc: 0.9759
Epoch 5/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1168 - acc: 0.9640 - val_loss: 0.0707 - val_acc: 0.9777
Epoch 6/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1066 - acc: 0.9676 - val_loss: 0.0752 - val_acc: 0.9759
Epoch 7/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1008 - acc: 0.9686 - val_loss: 0.0669 - val_acc: 0.9798
Epoch 8/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0923 - acc: 0.9714 - val_loss: 0.0657 - val_acc: 0.9792
Epoch 9/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0885 - acc: 0.9727 - val_loss: 0.0638 - val_acc: 0.9803
Epoch 10/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0840 - acc: 0.9740 - val_loss: 0.0603 - val_acc: 0.9810
Epoch 11/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0796 - acc: 0.9755 - val_loss: 0.0602 - val_acc: 0.9823
Epoch 12/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0712 - acc: 0.9778 - val_loss: 0.0559 - val_acc: 0.9822
```

```

val_loss: 0.0568 - val_acc: 0.9823
Epoch 13/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0698 - acc: 0.9779 -
val_loss: 0.0568 - val_acc: 0.9823
Epoch 14/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0696 - acc: 0.9779 -
val_loss: 0.0538 - val_acc: 0.9839
Epoch 15/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0665 - acc: 0.9790 -
val_loss: 0.0546 - val_acc: 0.9829
Epoch 16/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0591 - acc: 0.9815 -
val_loss: 0.0551 - val_acc: 0.9839
Epoch 17/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0566 - acc: 0.9820 -
val_loss: 0.0531 - val_acc: 0.9841
Epoch 18/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0580 - acc: 0.9813 -
val_loss: 0.0519 - val_acc: 0.9838
Epoch 19/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0545 - acc: 0.9833 -
val_loss: 0.0540 - val_acc: 0.9846
Epoch 20/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0539 - acc: 0.9822 -
val_loss: 0.0563 - val_acc: 0.9830

```

In [22]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

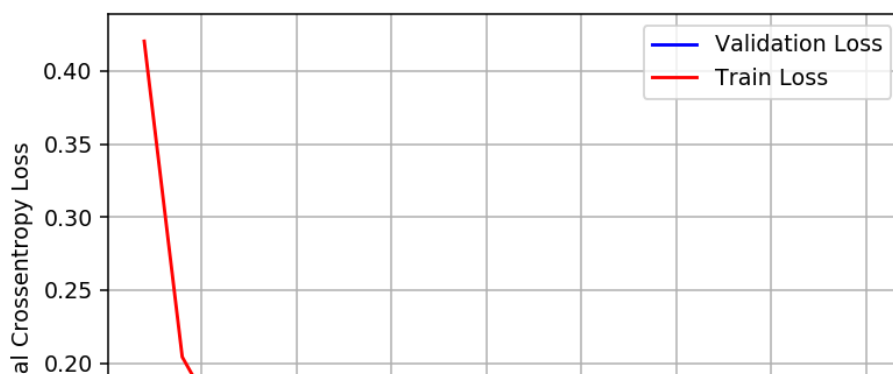
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

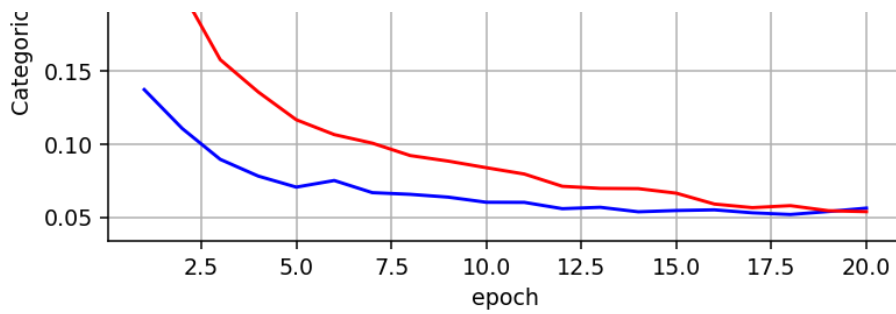
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05628199922235217

Test accuracy: 0.983





## 2. 3-hidden layer architecture (784-454-232-120-10)

### 2.1 MLP + ReLU activation + ADAM Optimizer

In [23]:

```
model_relu = Sequential()
model_relu.add(Dense(454, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(
seed=None)))
model_relu.add(Dense(232, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(120, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 454)	356390
dense_11 (Dense)	(None, 232)	105560
dense_12 (Dense)	(None, 120)	27960
dense_13 (Dense)	(None, 10)	1210
Total params: 491,120		
Trainable params: 491,120		
Non-trainable params: 0		

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 112us/step - loss: 0.2340 - acc: 0.9310 -
val_loss: 0.1052 - val_acc: 0.9669
Epoch 2/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0847 - acc: 0.9736 -
val_loss: 0.0826 - val_acc: 0.9733
Epoch 3/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0552 - acc: 0.9820 -
val_loss: 0.0708 - val_acc: 0.9772
Epoch 4/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0393 - acc: 0.9872 -
val_loss: 0.0727 - val_acc: 0.9781
Epoch 5/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0294 - acc: 0.9906 -
val_loss: 0.0772 - val_acc: 0.9790
Epoch 6/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0239 - acc: 0.9920 -
val_loss: 0.0793 - val_acc: 0.9768
Epoch 7/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0208 - acc: 0.9930 -
val_loss: 0.0834 - val_acc: 0.9787
Epoch 8/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0182 - acc: 0.9943 -
```

```

val_loss: 0.0749 - val_acc: 0.9806
Epoch 9/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0213 - acc: 0.9931 -
val_loss: 0.0836 - val_acc: 0.9803
Epoch 10/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0151 - acc: 0.9948 -
val_loss: 0.0828 - val_acc: 0.9806
Epoch 11/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0128 - acc: 0.9961 -
val_loss: 0.0906 - val_acc: 0.9783
Epoch 12/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0149 - acc: 0.9955 -
val_loss: 0.0765 - val_acc: 0.9828
Epoch 13/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0114 - acc: 0.9964 -
val_loss: 0.0977 - val_acc: 0.9783
Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0118 - acc: 0.9963 -
val_loss: 0.0844 - val_acc: 0.9806
Epoch 15/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0075 - acc: 0.9975 -
val_loss: 0.1109 - val_acc: 0.9792
Epoch 16/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0148 - acc: 0.9954 -
val_loss: 0.1126 - val_acc: 0.9785
Epoch 17/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0095 - acc: 0.9973 -
val_loss: 0.1064 - val_acc: 0.9791
Epoch 18/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0075 - acc: 0.9978 -
val_loss: 0.0876 - val_acc: 0.9820
Epoch 19/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0071 - acc: 0.9975 -
val_loss: 0.1041 - val_acc: 0.9798
Epoch 20/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0109 - acc: 0.9967 -
val_loss: 0.1178 - val_acc: 0.9778

```

In [24]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

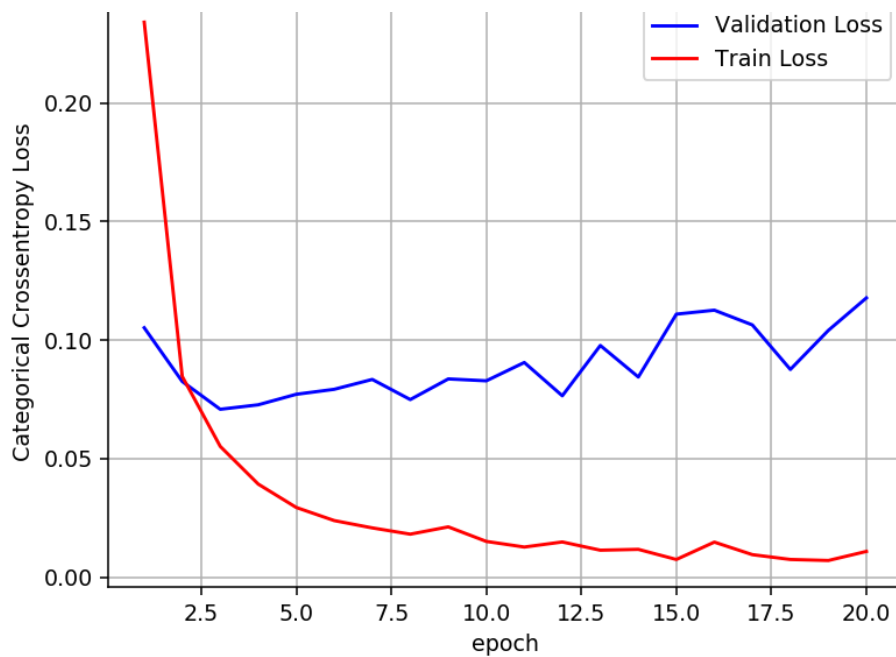
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.11779277188464211
Test accuracy: 0.9778

```



## 2.2 MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [25]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(454, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(232, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(120, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 454)	356390
batch_normalization_5 (Batch Normalization)	(None, 454)	1816
dense_15 (Dense)	(None, 232)	105560
batch_normalization_6 (Batch Normalization)	(None, 232)	928
dense_16 (Dense)	(None, 120)	27960
batch_normalization_7 (Batch Normalization)	(None, 120)	480
dense_17 (Dense)	(None, 10)	1210

```
=====
Total params: 494,344
Trainable params: 492,732
Non-trainable params: 1,612
=====
```

In [26]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 160us/step - loss: 0.1944 - acc: 0.9408 - val_loss: 0.1134 - val_acc: 0.9649
Epoch 2/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0745 - acc: 0.9767 - val_loss: 0.0873 - val_acc: 0.9727
Epoch 3/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0482 - acc: 0.9852 - val_loss: 0.0854 - val_acc: 0.9732
Epoch 4/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0408 - acc: 0.9872 - val_loss: 0.0816 - val_acc: 0.9747
Epoch 5/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0293 - acc: 0.9908 - val_loss: 0.0736 - val_acc: 0.9769
Epoch 6/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0257 - acc: 0.9919 - val_loss: 0.0774 - val_acc: 0.9754
Epoch 7/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0239 - acc: 0.9921 - val_loss: 0.0807 - val_acc: 0.9762
Epoch 8/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0215 - acc: 0.9929 - val_loss: 0.0791 - val_acc: 0.9771
Epoch 9/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0165 - acc: 0.9943 - val_loss: 0.0807 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0147 - acc: 0.9951 - val_loss: 0.0620 - val_acc: 0.9810
Epoch 11/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0157 - acc: 0.9944 - val_loss: 0.0745 - val_acc: 0.9798
Epoch 12/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0150 - acc: 0.9945 - val_loss: 0.0852 - val_acc: 0.9779
Epoch 13/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0117 - acc: 0.9962 - val_loss: 0.0717 - val_acc: 0.9817
Epoch 14/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0120 - acc: 0.9961 - val_loss: 0.0793 - val_acc: 0.9794
Epoch 15/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0126 - acc: 0.9959 - val_loss: 0.0717 - val_acc: 0.9817
Epoch 16/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0100 - acc: 0.9964 - val_loss: 0.0894 - val_acc: 0.9776
Epoch 17/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0103 - acc: 0.9964 - val_loss: 0.0794 - val_acc: 0.9809
Epoch 18/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0114 - acc: 0.9958 - val_loss: 0.0784 - val_acc: 0.9812
Epoch 19/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0085 - acc: 0.9973 - val_loss: 0.0833 - val_acc: 0.9814
Epoch 20/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0087 - acc: 0.9972 - val_loss: 0.0767 - val_acc: 0.9809
```



In [27]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

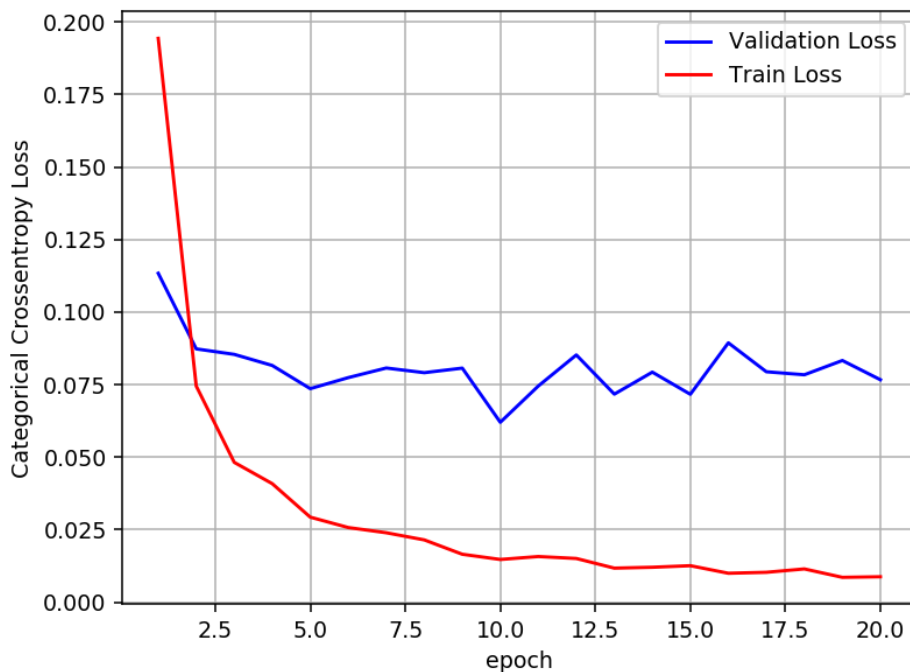
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07667597182277495

Test accuracy: 0.9809



## 2.3 MLP + Dropout + Adam Optimizer

In [28]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()
```

```

model_drop.add(Dense(454, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(
seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(232, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(120, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 454)	356390
batch_normalization_8 (Batch Normalization)	(None, 454)	1816
dropout_3 (Dropout)	(None, 454)	0
dense_19 (Dense)	(None, 232)	105560
batch_normalization_9 (Batch Normalization)	(None, 232)	928
dropout_4 (Dropout)	(None, 232)	0
dense_20 (Dense)	(None, 120)	27960
batch_normalization_10 (Batch Normalization)	(None, 120)	480
dropout_5 (Dropout)	(None, 120)	0
dense_21 (Dense)	(None, 10)	1210
Total params: 494,344		
Trainable params: 492,732		
Non-trainable params: 1,612		

In [29]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 193us/step - loss: 0.6018 - acc: 0.8176 - val_loss: 0.1744 - val_acc: 0.9470
Epoch 2/20
60000/60000 [=====] - 10s 162us/step - loss: 0.2632 - acc: 0.9222 - val_loss: 0.1235 - val_acc: 0.9624
Epoch 3/20
60000/60000 [=====] - 10s 162us/step - loss: 0.2041 - acc: 0.9395 - val_loss: 0.1069 - val_acc: 0.9677
Epoch 4/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1701 - acc: 0.9496 - val_loss: 0.0942 - val_acc: 0.9716
Epoch 5/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1459 - acc: 0.9572 - val_loss: 0.0813 - val_acc: 0.9748
Epoch 6/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1348 - acc: 0.9599 - val_loss: 0.0822 - val_acc: 0.9751
Epoch 7/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1255 - acc: 0.9629 - val_loss: 0.0822 - val_acc: 0.9751

```

```

oss: 0.0790 - val_acc: 0.9759
Epoch 8/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1144 - acc: 0.9666 - val_l
oss: 0.0781 - val_acc: 0.9756
Epoch 9/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1087 - acc: 0.9685 - val_l
oss: 0.0714 - val_acc: 0.9782
Epoch 10/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1014 - acc: 0.9693 - val_l
oss: 0.0708 - val_acc: 0.9787
Epoch 11/20
60000/60000 [=====] - 10s 158us/step - loss: 0.0937 - acc: 0.9721 - val_l
oss: 0.0685 - val_acc: 0.9814
Epoch 12/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0893 - acc: 0.9732 - val_l
oss: 0.0730 - val_acc: 0.9795
Epoch 13/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0903 - acc: 0.9730 - val_l
oss: 0.0662 - val_acc: 0.9811
Epoch 14/20
60000/60000 [=====] - 10s 162us/step - loss: 0.0806 - acc: 0.9759 - val_l
oss: 0.0627 - val_acc: 0.9824
Epoch 15/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0795 - acc: 0.9766 -
val_loss: 0.0629 - val_acc: 0.9812
Epoch 16/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0771 - acc: 0.9766 -
val_loss: 0.0590 - val_acc: 0.9829
Epoch 17/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0732 - acc: 0.9769 - val_l
oss: 0.0619 - val_acc: 0.9824
Epoch 18/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0721 - acc: 0.9777 - val_l
oss: 0.0598 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0686 - acc: 0.9784 - val_l
oss: 0.0582 - val_acc: 0.9821
Epoch 20/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0674 - acc: 0.9788 - val_l
oss: 0.0574 - val_acc: 0.9857

```

In [30]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

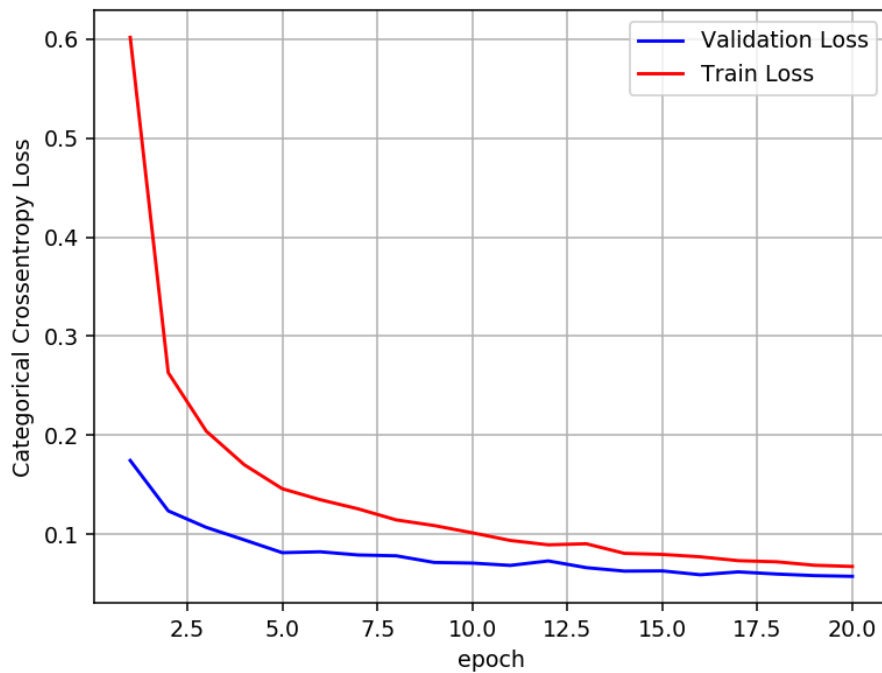
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.05740614543475094
Test accuracy: 0.9857

```



### 3. 5-hidden layer architecture (784-320-280-230-160-80-10)

#### 3.1 MLP + ReLU activation + ADAM Optimizer

In [31]:

```
model_relu = Sequential()
model_relu.add(Dense(320, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(
seed=None)))
model_relu.add(Dense(280, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(230, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(160, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(80, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 320)	251200
dense_23 (Dense)	(None, 280)	89880
dense_24 (Dense)	(None, 230)	64630
dense_25 (Dense)	(None, 160)	36960
dense_26 (Dense)	(None, 80)	12880
dense_27 (Dense)	(None, 10)	810

Total params: 456,360  
Trainable params: 456,360  
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 1/1 115/115 0.0000 0.0000 0.0000 0.0000

```

60000/60000 [=====] - /s 115us/step - loss: 0.2368 - acc: 0.9289 -
val_loss: 0.1155 - val_acc: 0.9640
Epoch 2/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0935 - acc: 0.9714 -
val_loss: 0.1012 - val_acc: 0.9694
Epoch 3/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0622 - acc: 0.9806 -
val_loss: 0.0916 - val_acc: 0.9725
Epoch 4/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0479 - acc: 0.9849 -
val_loss: 0.0869 - val_acc: 0.9759
Epoch 5/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0392 - acc: 0.9869 -
val_loss: 0.0801 - val_acc: 0.9761
Epoch 6/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0335 - acc: 0.9892 -
val_loss: 0.0940 - val_acc: 0.9729
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0280 - acc: 0.9914 -
val_loss: 0.0765 - val_acc: 0.9800
Epoch 8/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0259 - acc: 0.9914 -
val_loss: 0.0928 - val_acc: 0.9758
Epoch 9/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0235 - acc: 0.9924 -
val_loss: 0.0813 - val_acc: 0.9782
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0211 - acc: 0.9933 -
val_loss: 0.0806 - val_acc: 0.9801
Epoch 11/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0191 - acc: 0.9935 -
val_loss: 0.1081 - val_acc: 0.9749
Epoch 12/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0183 - acc: 0.9944 -
val_loss: 0.0765 - val_acc: 0.9807
Epoch 13/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0142 - acc: 0.9956 -
val_loss: 0.0975 - val_acc: 0.9775
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0187 - acc: 0.9944 -
val_loss: 0.0867 - val_acc: 0.9782
Epoch 15/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0182 - acc: 0.9946 -
val_loss: 0.0775 - val_acc: 0.9813
Epoch 16/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0105 - acc: 0.9968 -
val_loss: 0.0994 - val_acc: 0.9798
Epoch 17/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0155 - acc: 0.9950 -
val_loss: 0.1075 - val_acc: 0.9779
Epoch 18/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0135 - acc: 0.9959 -
val_loss: 0.0959 - val_acc: 0.9795
Epoch 19/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0123 - acc: 0.9963 -
val_loss: 0.0997 - val_acc: 0.9781
Epoch 20/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0072 - acc: 0.9979 -
val_loss: 0.0924 - val_acc: 0.9812

```

In [32]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

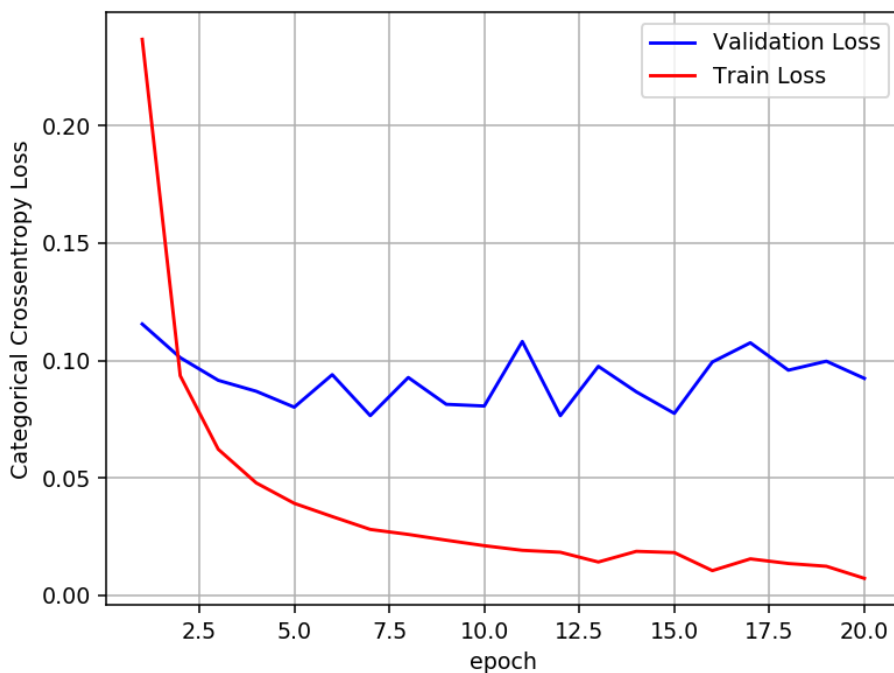
```

```
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09238322618362185  
Test accuracy: 0.9812



## 3.2 MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [33]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(320, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(280, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(230, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())
```

```

model_batch.add(Dense(160, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(80, activation='relu', kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 320)	251200
batch_normalization_11 (Batch Normalization)	(None, 320)	1280
dense_29 (Dense)	(None, 280)	89880
batch_normalization_12 (Batch Normalization)	(None, 280)	1120
dense_30 (Dense)	(None, 230)	64630
batch_normalization_13 (Batch Normalization)	(None, 230)	920
dense_31 (Dense)	(None, 160)	36960
batch_normalization_14 (Batch Normalization)	(None, 160)	640
dense_32 (Dense)	(None, 80)	12880
batch_normalization_15 (Batch Normalization)	(None, 80)	320
dense_33 (Dense)	(None, 10)	810
Total params: 460,640		
Trainable params: 458,500		
Non-trainable params: 2,140		

In [34]:

```

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 201us/step - loss: 0.2278 - acc: 0.9324 - val_loss: 0.1180 - val_acc: 0.9658
Epoch 2/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0869 - acc: 0.9731 - val_loss: 0.0921 - val_acc: 0.9710
Epoch 3/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0596 - acc: 0.9808 - val_loss: 0.0831 - val_acc: 0.9722
Epoch 4/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0482 - acc: 0.9842 - val_loss: 0.0878 - val_acc: 0.9730
Epoch 5/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0403 - acc: 0.9863 - val_loss: 0.0907 - val_acc: 0.9729
Epoch 6/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0339 - acc: 0.9892 - val_loss: 0.0968 - val_acc: 0.9739
Epoch 7/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0306 - acc: 0.9895 - val_loss: 0.0820 - val_acc: 0.9774
Epoch 8/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0265 - acc: 0.9913 - val_loss: 0.0915 - val_acc: 0.9760
Epoch 9/20
60000/60000 [=====] - 9s 126us/step - loss: 0.0274 - acc: 0.9900

```

```

60000/60000 [=====] - 8s 126us/step - loss: 0.0274 - acc: 0.9909 -
val_loss: 0.0776 - val_acc: 0.9779
Epoch 10/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0221 - acc: 0.9927 -
val_loss: 0.0723 - val_acc: 0.9795
Epoch 11/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0227 - acc: 0.9925 -
val_loss: 0.0782 - val_acc: 0.9791
Epoch 12/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0180 - acc: 0.9937 -
val_loss: 0.0942 - val_acc: 0.9752
Epoch 13/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0196 - acc: 0.9936 -
val_loss: 0.0708 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0168 - acc: 0.9943 -
val_loss: 0.0684 - val_acc: 0.9825
Epoch 15/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0163 - acc: 0.9945 -
val_loss: 0.0819 - val_acc: 0.9794
Epoch 16/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0129 - acc: 0.9959 -
val_loss: 0.0902 - val_acc: 0.9789
Epoch 17/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0184 - acc: 0.9938 -
val_loss: 0.0921 - val_acc: 0.9765
Epoch 18/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0139 - acc: 0.9953 -
val_loss: 0.0880 - val_acc: 0.9790
Epoch 19/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0143 - acc: 0.9949 -
val_loss: 0.0859 - val_acc: 0.9806
Epoch 20/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0123 - acc: 0.9960 -
val_loss: 0.0813 - val_acc: 0.9796

```

In [35]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

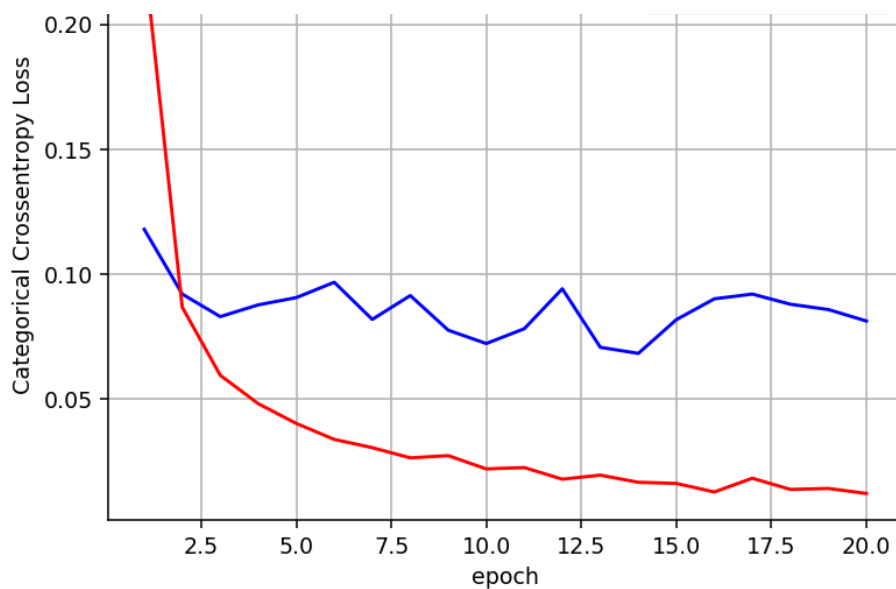
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0813284714843132  
Test accuracy: 0.9796







### 3.3 MLP + Dropout + Adam Optimizer

In [36]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(320, activation='relu', input_shape=(input_dim,)), kernel_initializer=he_normal(
(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(280, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(230, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(160, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(80, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 320)	251200
batch_normalization_16 (Batch Normalization)	(None, 320)	1280
dropout_6 (Dropout)	(None, 320)	0
dense_35 (Dense)	(None, 280)	89880
batch_normalization_17 (Batch Normalization)	(None, 280)	1120
dropout_7 (Dropout)	(None, 280)	0

dense_36 (Dense)	(None, 230)	64630
batch_normalization_18 (Batch Normalization)	(None, 230)	920
dropout_8 (Dropout)	(None, 230)	0
dense_37 (Dense)	(None, 160)	36960
batch_normalization_19 (Batch Normalization)	(None, 160)	640
dropout_9 (Dropout)	(None, 160)	0
dense_38 (Dense)	(None, 80)	12880
batch_normalization_20 (Batch Normalization)	(None, 80)	320
dropout_10 (Dropout)	(None, 80)	0
dense_39 (Dense)	(None, 10)	810
=====		
Total params: 460,640		
Trainable params: 458,500		
Non-trainable params: 2,140		

In [37]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 17s 280us/step - loss: 1.1816 - acc: 0.6242 - val_loss: 0.2850 - val_acc: 0.9147
Epoch 2/20
60000/60000 [=====] - 12s 204us/step - loss: 0.4393 - acc: 0.8740 - val_loss: 0.1872 - val_acc: 0.9476
Epoch 3/20
60000/60000 [=====] - 12s 208us/step - loss: 0.3221 - acc: 0.9100 - val_loss: 0.1546 - val_acc: 0.9544
Epoch 4/20
60000/60000 [=====] - 12s 203us/step - loss: 0.2687 - acc: 0.9280 - val_loss: 0.1426 - val_acc: 0.9600
Epoch 5/20
60000/60000 [=====] - 12s 207us/step - loss: 0.2335 - acc: 0.9384 - val_loss: 0.1239 - val_acc: 0.9668
Epoch 6/20
60000/60000 [=====] - 12s 204us/step - loss: 0.2128 - acc: 0.9434 - val_loss: 0.1096 - val_acc: 0.9716
Epoch 7/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1909 - acc: 0.9496 - val_loss: 0.1030 - val_acc: 0.9736
Epoch 8/20
60000/60000 [=====] - 12s 204us/step - loss: 0.1797 - acc: 0.9533 - val_loss: 0.0979 - val_acc: 0.9743
Epoch 9/20
60000/60000 [=====] - 12s 202us/step - loss: 0.1680 - acc: 0.9556 - val_loss: 0.0988 - val_acc: 0.9745
Epoch 10/20
60000/60000 [=====] - 12s 202us/step - loss: 0.1566 - acc: 0.9589 - val_loss: 0.0965 - val_acc: 0.9741
Epoch 11/20
60000/60000 [=====] - 12s 203us/step - loss: 0.1488 - acc: 0.9612 - val_loss: 0.0908 - val_acc: 0.9772
Epoch 12/20
60000/60000 [=====] - 12s 203us/step - loss: 0.1420 - acc: 0.9620 - val_loss: 0.0880 - val_acc: 0.9781
Epoch 13/20
60000/60000 [=====] - 12s 204us/step - loss: 0.1372 - acc: 0.9635 - val_loss: 0.0856 - val_acc: 0.9778
Epoch 14/20
60000/60000 [=====] - 12s 199us/step - loss: 0.1392 - acc: 0.9639 - val_loss: 0.0850 - val_acc: 0.9781
```

```

oss: 0.0859 - val_acc: 0.9781
Epoch 15/20
60000/60000 [=====] - 12s 198us/step - loss: 0.1269 - acc: 0.9662 - val_l
oss: 0.0856 - val_acc: 0.9771
Epoch 16/20
60000/60000 [=====] - 12s 197us/step - loss: 0.1250 - acc: 0.9675 - val_l
oss: 0.0855 - val_acc: 0.9792
Epoch 17/20
60000/60000 [=====] - 12s 199us/step - loss: 0.1213 - acc: 0.9679 - val_l
oss: 0.0823 - val_acc: 0.9797
Epoch 18/20
60000/60000 [=====] - 12s 199us/step - loss: 0.1129 - acc: 0.9692 - val_l
oss: 0.0795 - val_acc: 0.9792
Epoch 19/20
60000/60000 [=====] - 12s 199us/step - loss: 0.1123 - acc: 0.9700 - val_l
oss: 0.0824 - val_acc: 0.9773
Epoch 20/20
60000/60000 [=====] - 12s 199us/step - loss: 0.1088 - acc: 0.9710 - val_l
oss: 0.0758 - val_acc: 0.9806

```

In [38]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

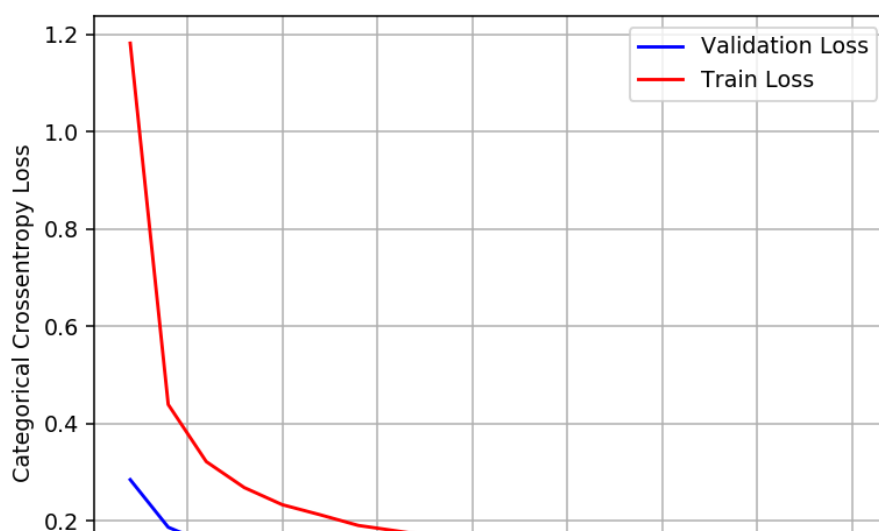
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

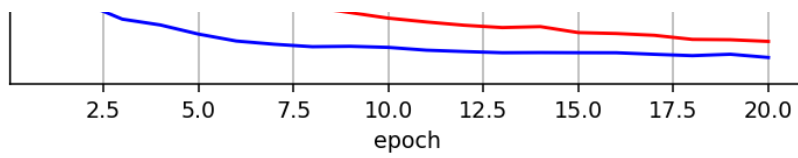
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07577368428185582  
Test accuracy: 0.9806





## 4. CONCLUSION

In [40]:

```
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Numer of Layers", "BN", "Dropout", "Accuracy"]
x.add_row(["2", "NO", "NO", 0.981])
x.add_row(["2", "YES", "NO", 0.978])
x.add_row(["2", "NO", 0.5, 0.983])

x.add_row(["3", "NO", "NO", 0.977])
x.add_row(["3", "YES", "NO", 0.980])
x.add_row(["3", "NO", 0.5, 0.985])

x.add_row(["5", "NO", "NO", 0.981])
x.add_row(["5", "YES", "NO", 0.979])
x.add_row(["5", "NO", 0.5, 0.980])
print(x)
```

Numer of Layers	BN	Dropout	Accuracy
2	NO	NO	0.981
2	YES	NO	0.978
2	NO	0.5	0.983
3	NO	NO	0.977
3	YES	NO	0.98
3	NO	0.5	0.985
5	NO	NO	0.981
5	YES	NO	0.979
5	NO	0.5	0.98

In [ ]: